

#DLUPC

Day 1 Lecture 4

Multilayer Perceptron MLP



Xavier Giro-i-Nieto

xavier.giro@upc.edu

Associate Professor

Universitat Politècnica de Catalunya
Technical University of Catalonia



Acknowledgements



Kevin McGuinness

kevin.mcguinness@dcu.ie

Research Fellow

Insight Centre for Data Analytics
Dublin City University



Eva Mohedano

eva.mohedano@insight-centre.org

Postdoctoral Researcher

Insight Centre for Data Analytics
Dublin City University



Elisa Sayrol

elisa.sayrol@upc.edu

Associate Professor

ETSETB TelecomBCN
Universitat Politècnica de Catalunya



Antonio Bonafonte

antonio.bonafonte@upc.edu

Associate Professor

ETSETB TelecomBCN
Universitat Politècnica de Catalunya

Video lecture

DEEP LEARNING FOR ARTIFICIAL INTELLIGENCE

Master Course UPC ETSETB TelecomBCN Barcelona, Autumn 2017



Instructors



Organizers



Supporters



+ info: <http://dlai.deeplearning.barcelona>

[\[course site\]](#)

Day 2 Lecture 1

Multilayer Perceptron



Elisa Sayrol



UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH
Department of Signal Theory
and Communications
Image Processing Group

Overview

- Limitations the perceptron model
- Multilayer perceptron

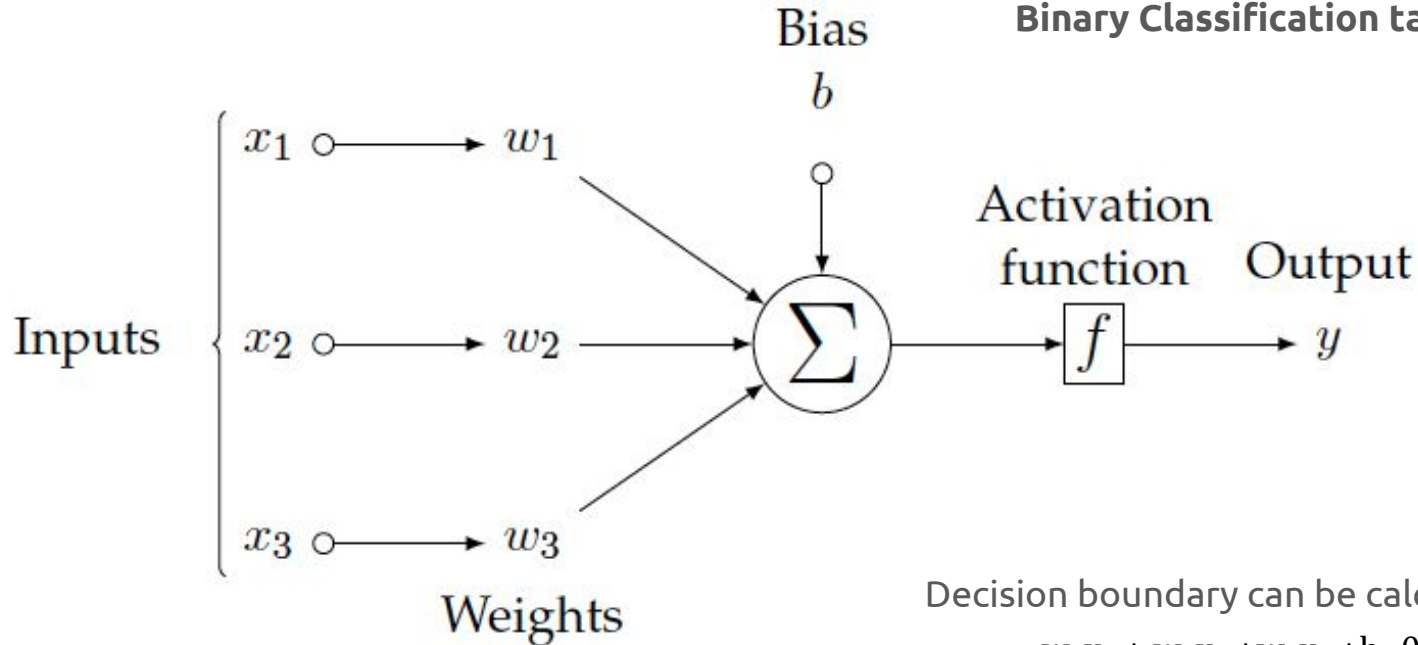
Overview

- **Limitations the perceptron model**
- Multilayer perceptron

Single Neuron

If the weighted sum of the input exceeds a threshold the neuron fires a signal.

Binary Classification task

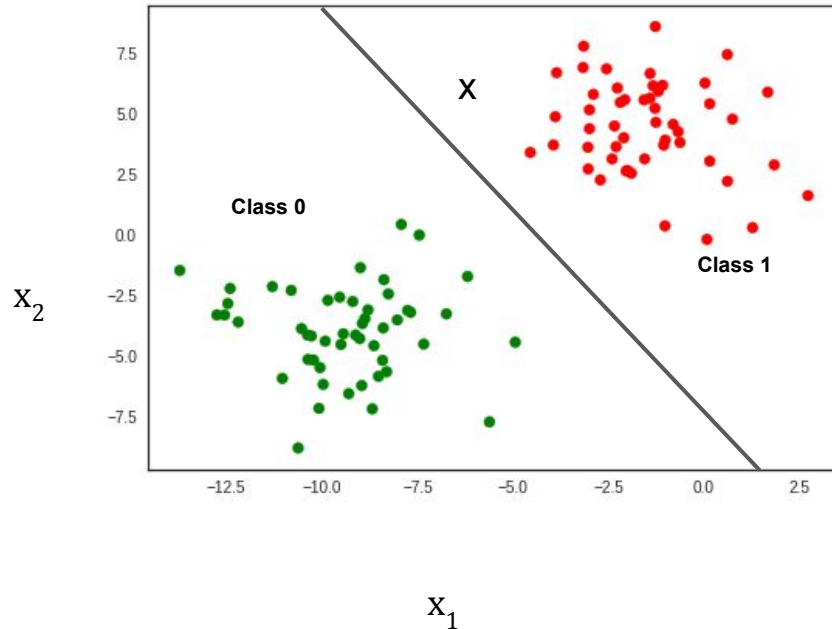


Decision boundary can be calculated by:

$$w_1x_1 + w_2x_2 + w_3x_3 + b = 0$$

Linear decision boundary

2D input space data



$$f(x) = \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

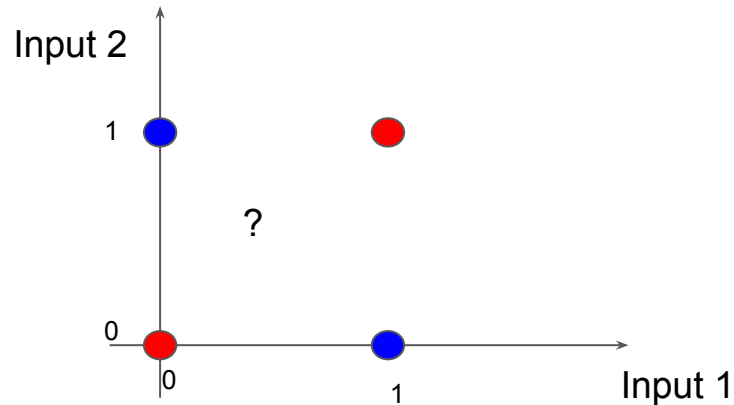
The XOR problem

Limitation: Data might be **non linearly separable**

→ One single neuron is not enough

XOR logic table

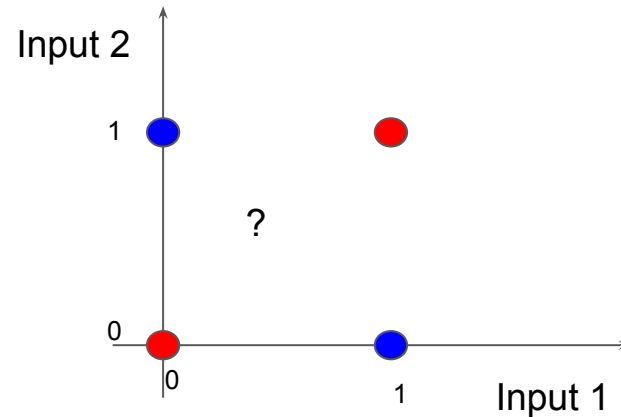
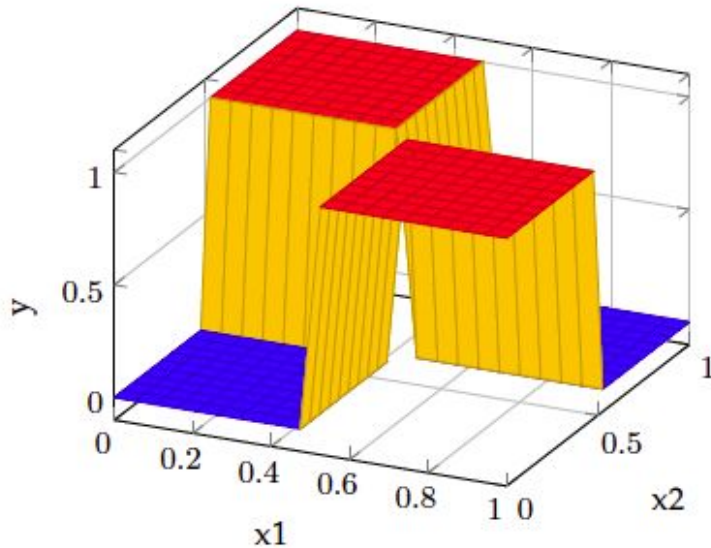
Input 1	Input 2	Desired Output
0	0	0
0	1	1
1	0	1
1	1	0



The XOR problem

Limitation: Data might be **non linearly separable**

→ One single neuron is not enough



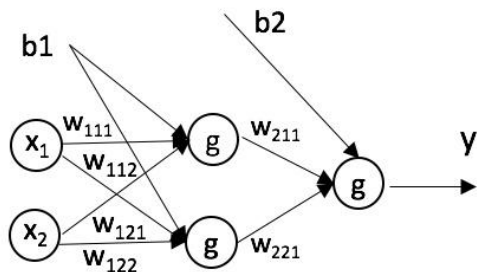
The limitations of the perceptron



Minsky, Marvin, and Seymour A. Papert. [Perceptrons: An introduction to computational geometry](#). 1969

The XOR problem

Given the following network to obtain a XNOR operation, Indicate which set of parameters are correct (a, b, c or d):

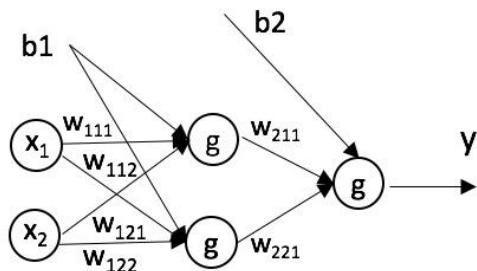


Input vector (x_1, x_2)	Class XNOR
(0,0)	1
(0,1)	0
(1,0)	0
(1,1)	1

	w_{111}	w_{112}	w_{121}	w_{122}	b_1	w_{211}	w_{221}	b_2
a	2	2	2	2	-1	2	2	-1
b	-2	2	2	-2	-1	2	2	1
c	-2	2	2	-2	-1	2	-2	1
d	-2	2	2	-2	-1	2	-2	-1

The XOR problem

Given the following network to obtain a XNOR operation, Indicate which set of parameters are correct (a, b, c or d):

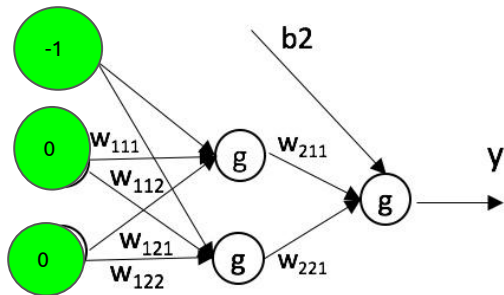


Input vector (x_1, x_2)	Class XNOR
(0,0)	1
(0,1)	0
(1,0)	0
(1,1)	1

	w_{111}	w_{112}	w_{121}	w_{122}	b_1	w_{211}	w_{221}	b_2
a	2	2	2	2	-1	2	2	-1
b	-2	2	2	-2	-1	2	2	1
c	-2	2	2	-2	-1	2	-2	1
d	-2	2	2	-2	-1	2	-2	-1

The XOR problem

Given the following network to obtain a XNOR operation, Indicate which set of parameters are correct (a, b, c or d):

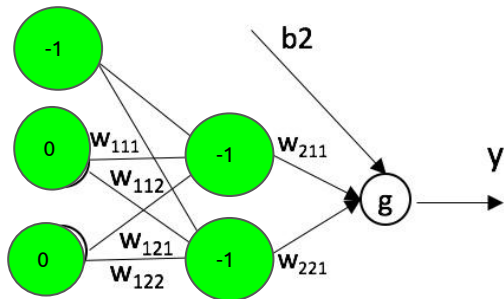


Input vector (x1,x2)	Class XNOR
(0,0)	1
(0,1)	0
(1,0)	0
(1,1)	1

	w_{111}	w_{112}	w_{121}	w_{122}	b_1	w_{211}	w_{221}	b_2
a	2	2	2	2	-1	2	2	-1
b	-2	2	2	-2	-1	2	2	1
c	-2	2	2	-2	-1	2	-2	1
d	-2	2	2	-2	-1	2	-2	-1

The XOR problem

Given the following network to obtain a XNOR operation, Indicate which set of parameters are correct (a, b, c or d):

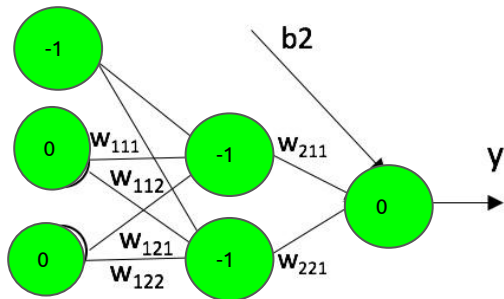


Input vector (x1,x2)	Class XNOR
(0,0)	1
(0,1)	0
(1,0)	0
(1,1)	1

	w_{111}	w_{112}	w_{121}	w_{122}	b_1	w_{211}	w_{221}	b_2
a	2	2	2	2	-1	2	2	-1
b	-2	2	2	-2	-1	2	2	1
c	-2	2	2	-2	-1	2	-2	1
d	-2	2	2	-2	-1	2	-2	-1

The XOR problem

Given the following network to obtain a XNOR operation, Indicate which set of parameters are correct (a, b, c or d):

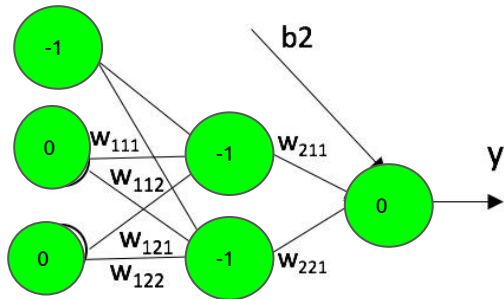


Input vector (x1,x2)	Class XNOR
(0,0)	1
(0,1)	0
(1,0)	0
(1,1)	1

	w_{111}	w_{112}	w_{121}	w_{122}	b_1	w_{211}	w_{221}	b_2
a	2	2	2	2	-1	2	2	-1
b	-2	2	2	-2	-1	2	2	1
c	-2	2	2	-2	-1	2	-2	1
d	-2	2	2	-2	-1	2	2	-1

The XOR problem

Given the following network to obtain a XNOR operation, Indicate which set of parameters are correct (a, b, c or d):

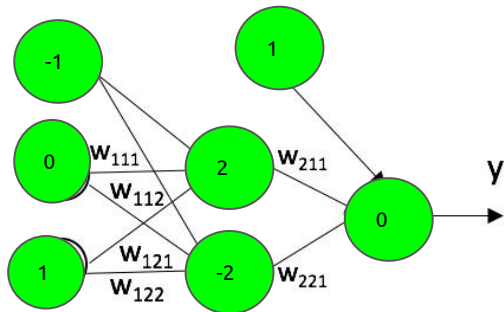


Input vector (x1,x2)	Class XNOR
(0,0)	1
(0,1)	0
(1,0)	0
(1,1)	1

	w_{111}	w_{112}	w_{121}	w_{122}	b_1	w_{211}	w_{221}	b_2
a	2	2	2	2	-1	2	2	-1
b	-2	2	2	-2	-1	2	2	1
c	-2	2	2	-2	-1	2	-2	1
d	-2	2	2	-2	-1	2	-2	-1

The XOR problem

Given the following network to obtain a XNOR operation, Indicate which set of parameters are correct (a, b, c or d):

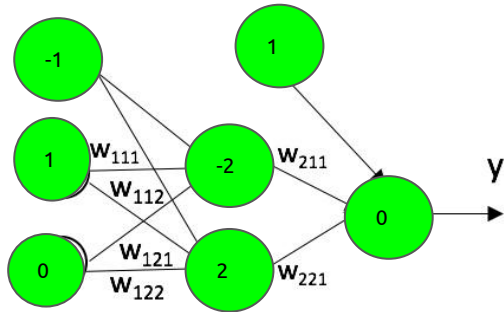


Input vector (x1,x2)	Class XNOR
(0,0)	1
(0,1)	0
(1,0)	0
(1,1)	1

	w_{111}	w_{112}	w_{121}	w_{122}	b_1	w_{211}	w_{221}	b_2
a	2	2	2	2	-1	2	2	-1
b	-2	2	2	-2	-1	2	2	1
c	-2	2	2	-2	-1	2	-2	1
d	-2	2	2	-2	-1	2	-2	-1

The XOR problem

Given the following network to obtain a XNOR operation, Indicate which set of parameters are correct (a, b, c or d):

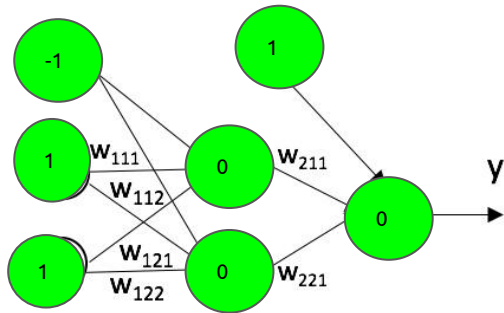


Input vector (x1,x2)	Class XNOR
(0,0)	1
(0,1)	0
(1,0)	0
(1,1)	1

	w_{111}	w_{112}	w_{121}	w_{122}	b_1	w_{211}	w_{221}	b_2
a	2	2	2	2	-1	2	2	-1
b	-2	2	2	-2	-1	2	2	1
c	-2	2	2	-2	-1	2	-2	1
d	-2	2	2	-2	-1	2	-2	-1

The XOR problem

Given the following network to obtain a XNOR operation, Indicate which set of parameters are correct (a, b, c or d):



Input vector (x1,x2)	Class XNOR
(0,0)	1
(0,1)	0
(1,0)	0
(1,1)	1

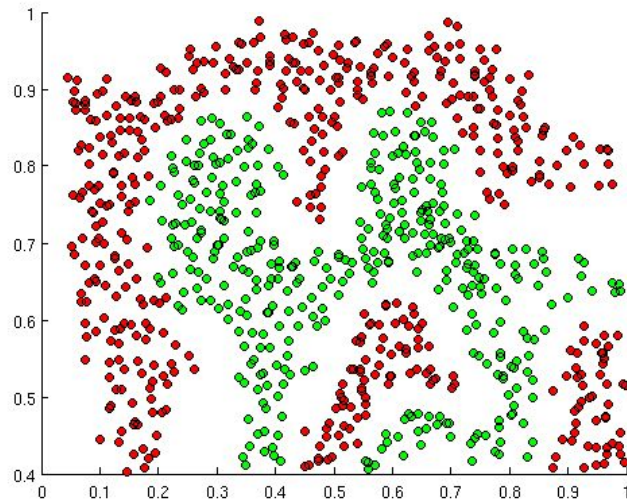
	w_{111}	w_{112}	w_{121}	w_{122}	b_1	w_{211}	w_{221}	b_2
a	2	2	2	2	-1	2	2	-1
b	-2	2	2	-2	-1	2	2	1
c	-2	2	2	-2	-1	2	-2	1
d	-2	2	2	-2	-1	2	-2	-1

Non-linear decision boundaries

Linear models can only produce linear decision boundaries

Real world data often needs a non-linear decision boundary

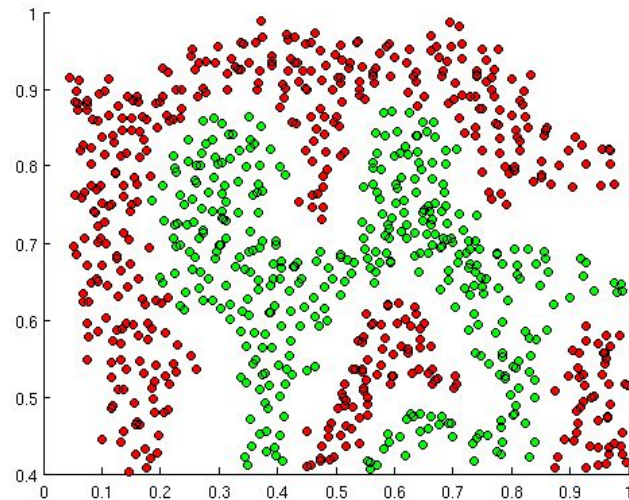
- Images
- Audio
- Text



Non-linear decision boundaries

What can we do?

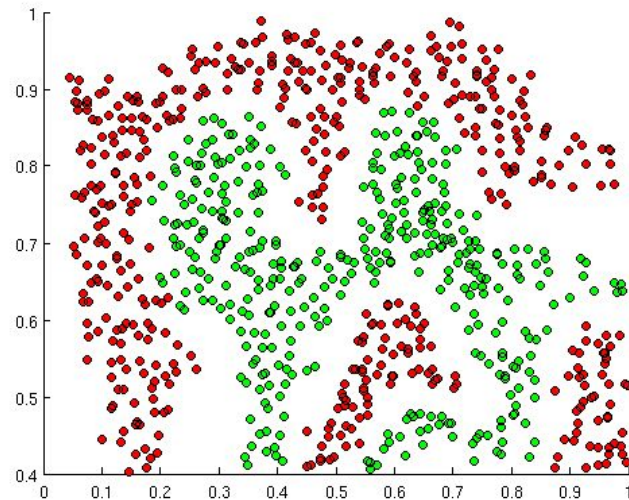
1. Use a non-linear classifier
 - Decision trees (and forests)
 - K nearest neighbours
2. Engineer a suitable representation
 - One in which features are more linearly separable
 - Then use a linear model
3. Engineer a kernel
 - Design a kernel $K(\mathbf{x}_1, \mathbf{x}_2)$
 - Use kernel methods (e.g. SVM)
4. Learn a suitable representation space from the data
 - Deep learning, deep neural networks
 - Boosted cascade classifiers like Viola Jones also take this approach



Non-linear decision boundaries

What can we do?

1. Use a non-linear classifier
 - Decision trees (and forests)
 - K nearest neighbours
2. Engineer a suitable representation
 - One in which features are more linearly separable
 - Then use a linear model
3. Engineer a kernel
 - Design a kernel $K(\mathbf{x}_1, \mathbf{x}_2)$
 - Use kernel methods (e.g. SVM)
4. Learn a suitable representation space from the data
 - **Deep learning, deep neural networks**
 - Boosted cascade classifiers like Viola Jones also take this approach



Overview

- Limitations the perceptron model
- **Multilayer perceptron**

Multilayer perceptrons

When each node in each layer is a linear combination of **all inputs from the previous layer** then the network is called a multilayer perceptron (MLP)

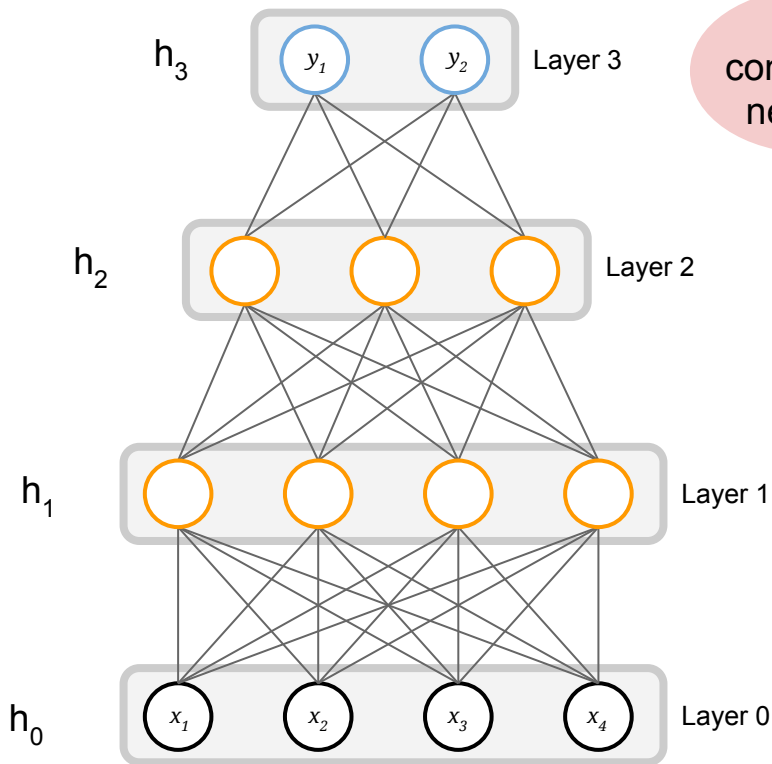
Weights can be organized into matrices.

Forward pass computes

$$\mathbf{h}_0 = \mathbf{x}$$

$$\mathbf{h}^{(t)} = g(W^{(t)}\mathbf{h}^{(t-1)} + \mathbf{b}^{(t)})$$

$$f(\mathbf{x}) = \mathbf{h}^{(L)}$$



Multilayer perceptrons

W_1				h_0	b_1
w_{11}	w_{12}	w_{13}	w_{14}		
w_{21}	w_{22}	w_{23}	w_{24}		
w_{31}	w_{32}	w_{33}	w_{34}		
w_{41}	w_{42}	w_{43}	w_{44}		
x_1				x_2	b_1
				x_3	b_2
				x_4	b_3
					b_4

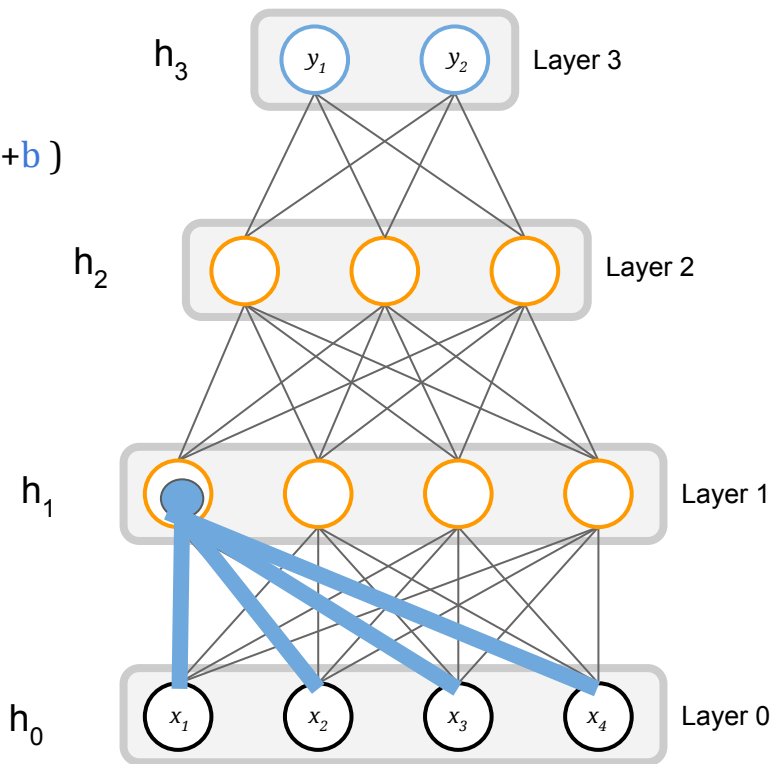
$$h_{11} = g(\mathbf{w}\mathbf{x} + \mathbf{b})$$

Forward pass computes

$$\mathbf{h}_0 = \mathbf{x}$$

$$\mathbf{h}^{(t)} = g(W^{(t)}\mathbf{h}^{(t-1)} + \mathbf{b}^{(t)})$$

$$f(\mathbf{x}) = \mathbf{h}^{(L)}$$



Multilayer perceptrons

W_1				h_0	b_1
w_{11}	w_{12}	w_{13}	w_{14}	x_1	b_1
w_{21}	w_{22}	w_{23}	w_{24}	x_2	b_2
w_{31}	w_{32}	w_{33}	w_{34}	x_3	b_3
w_{41}	w_{42}	w_{43}	w_{44}	x_4	b_4

$$h_{11} = g(\mathbf{w}\mathbf{x} + \mathbf{b})$$

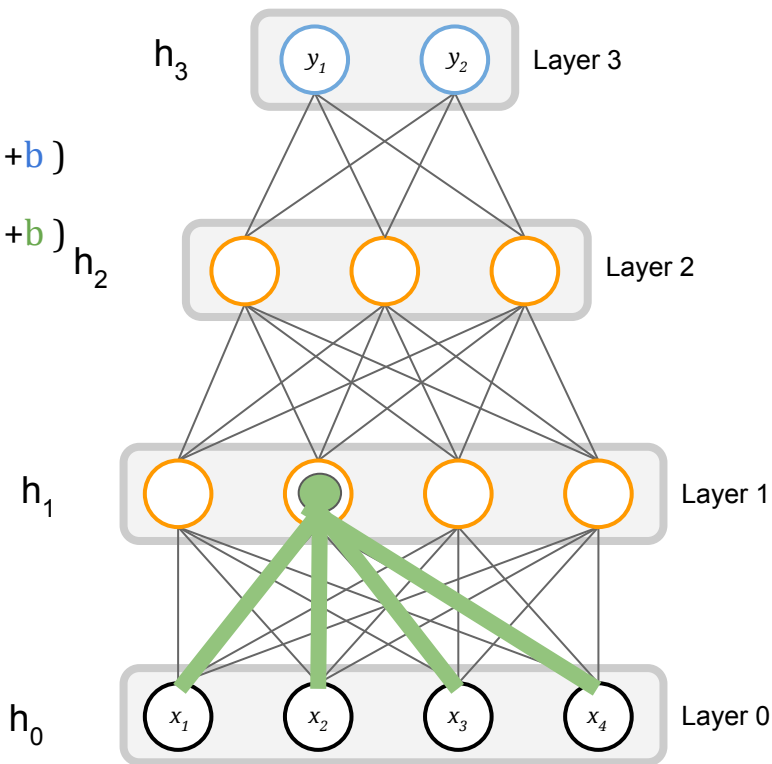
$$h_{12} = g(\mathbf{w}\mathbf{x} + \mathbf{b})$$

Forward pass computes

$$\mathbf{h}_0 = \mathbf{x}$$

$$\mathbf{h}^{(t)} = g(W^{(t)}\mathbf{h}^{(t-1)} + \mathbf{b}^{(t)})$$

$$f(\mathbf{x}) = \mathbf{h}^{(L)}$$



Demo: MNIST digit classification

MNIST

- Popular dataset of handwritten digits
- 60,000 training examples
- 10,000 test examples
- 10 classes (digits 0-9)
- <http://yann.lecun.com/exdb/mnist/>
- 28x28 grayscale images (784D)

Objective

- Learn a function $y = f(x)$ that predicts the digit from the image
- Measure accuracy on test set

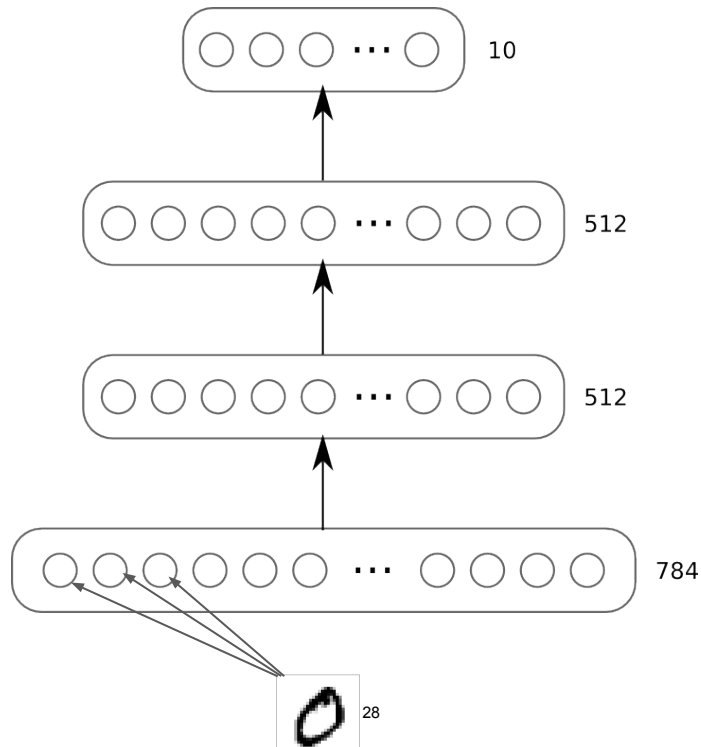


Multilayer perceptrons - MLPs

How many parameters contains the following MLP ?

Model

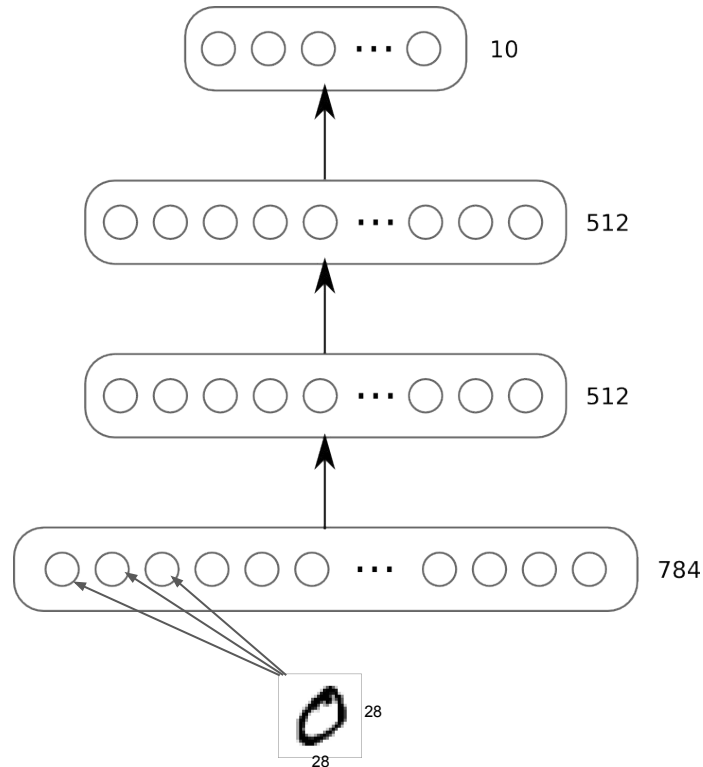
- 3 layer neural network (2 hidden layers)
- Tanh units (activation function)
- 512-512-10
- **Softmax** on top layer



Multilayer perceptrons - MLPs

How many parameters contains the following MLP ?

Layer	#Weights	#Biases	Total
1	784 x 512	512	401,920
2	512 x 512	512	262,656
3	512 x 10	10	5,130
			669,706



Multilayer perceptrons - MLPs

Which of the statements below is true?

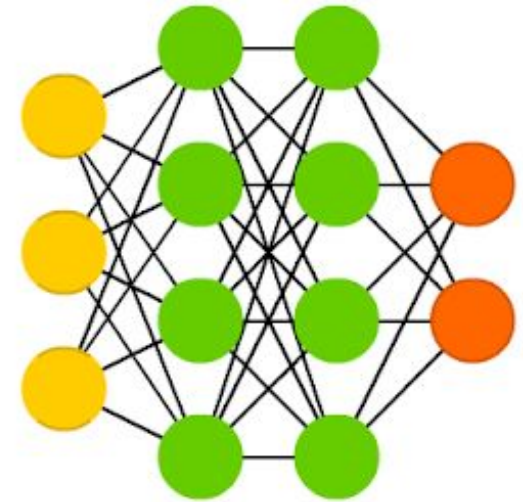
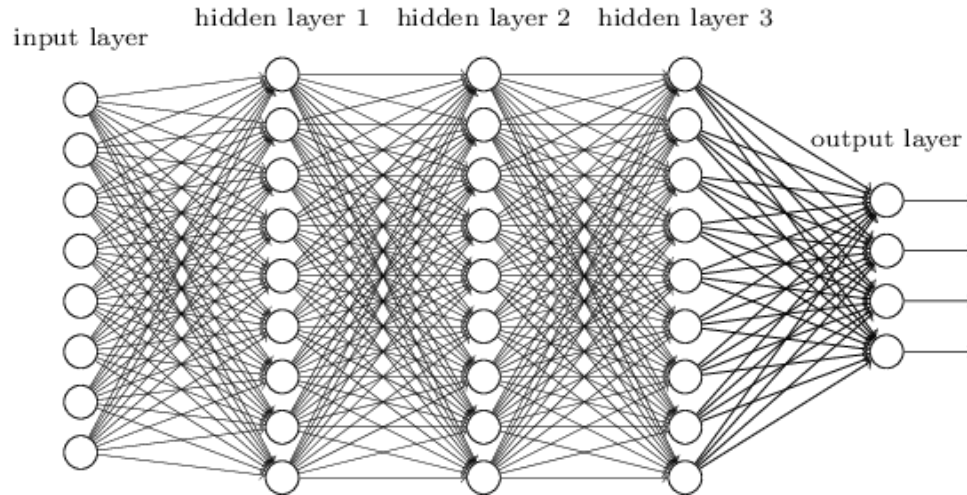
- A. Multi-layer perceptrons and CNNs are the same kind of networks
- B. Each node in a given layer is connected to all inputs from the previous layer
- C. In multi-layer perceptrons, the hidden layers are connected only to the input layer
- D. There are no hidden layers in the Multi-layer perceptron only in deep networks

Multilayer perceptrons - MLPs

Which of the statements below is true?

- A. Multi-layer perceptrons and CNNs are the same kind of networks
- B. Each node in a given layer is connected to all inputs from the previous layer**
- C. In multi-layer perceptrons, the hidden layers are connected only to the input layer
- D. There are no hidden layers in the Multi-layer perceptron only in deep networks

Deep Neural Networks (DNN)



- Input Cell
- Hidden Cell
- Output Cell

Universal approximation theorem

Universal approximation theorem states that “the standard multilayer feed-forward network with a **single hidden layer**, which contains **finite number of hidden neurons**, is a **universal approximator** among continuous functions on compact subsets of \mathbb{R}^n , under mild assumptions on the activation function.”

If a 2 layer NN is a universal approximator, then why do we need **deep nets** ??

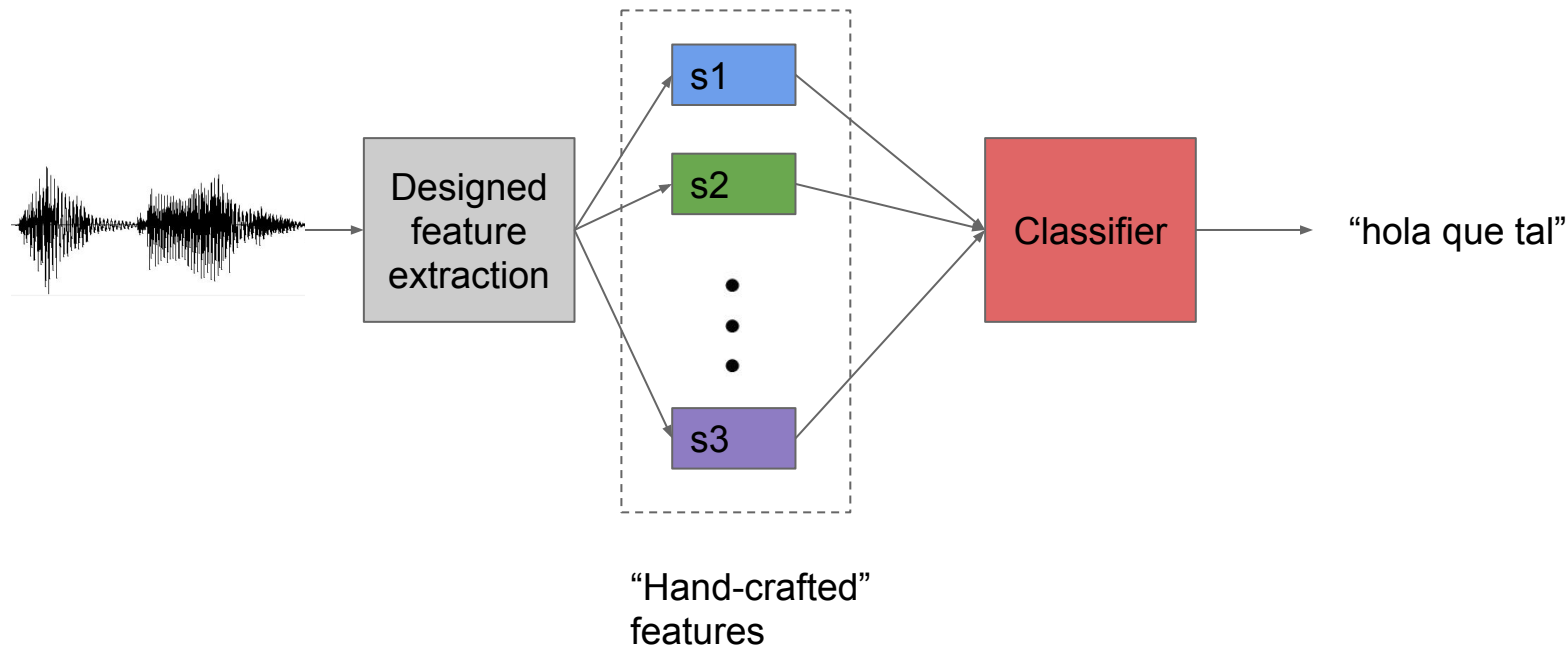
The universal approximation theorem:

- Says nothing about the how easy/difficult it is to fit such approximators
- Needs a “finite number of hidden neurons”: finite may be extremely large

In practice, deep nets can usually represent more complex functions with less total neurons (and therefore, less parameters)

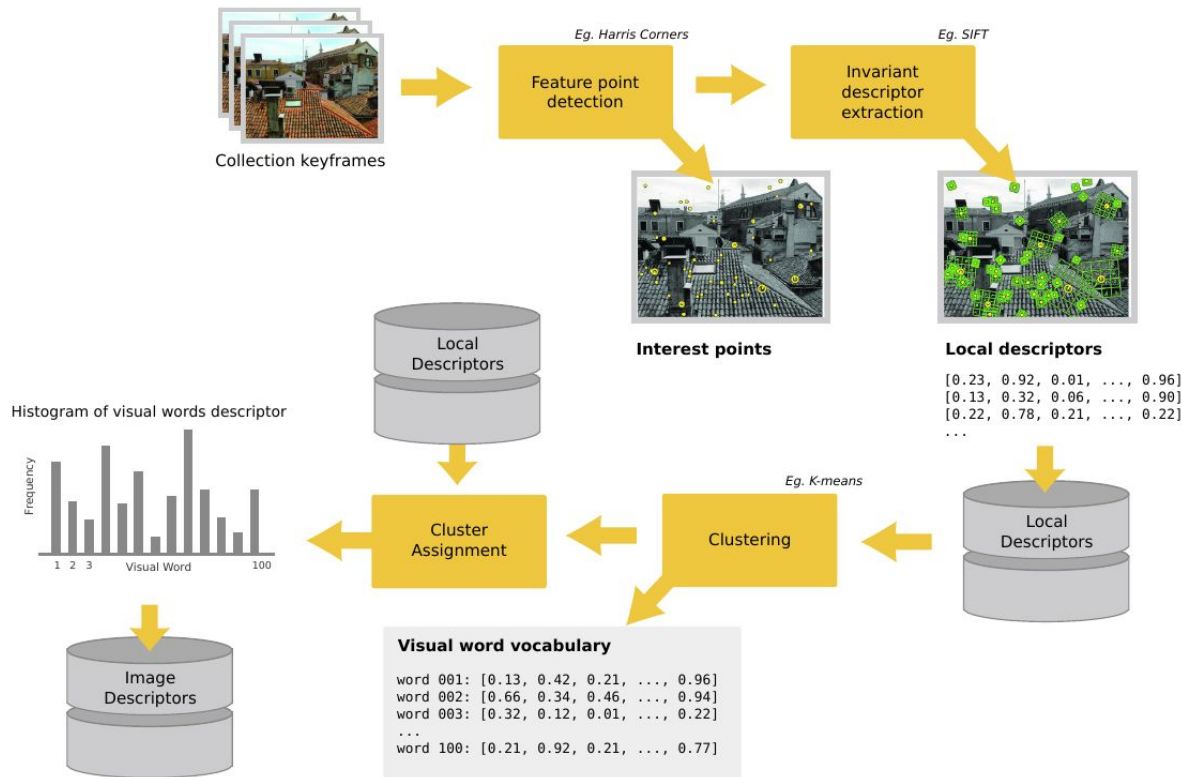
Classic Machine Learning...

Feature engineering for Automatic Speech Recognition (ASR).



Classic Machine Learning...

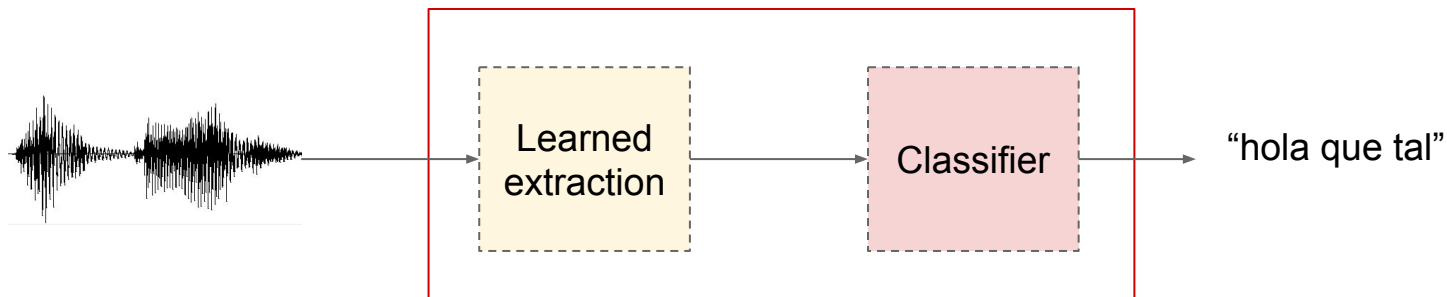
Feature engineering for Computer Vision.



Classic Machine Learning vs Deep Learning

Learn the representations as well, not only the final mapping → **end2end**

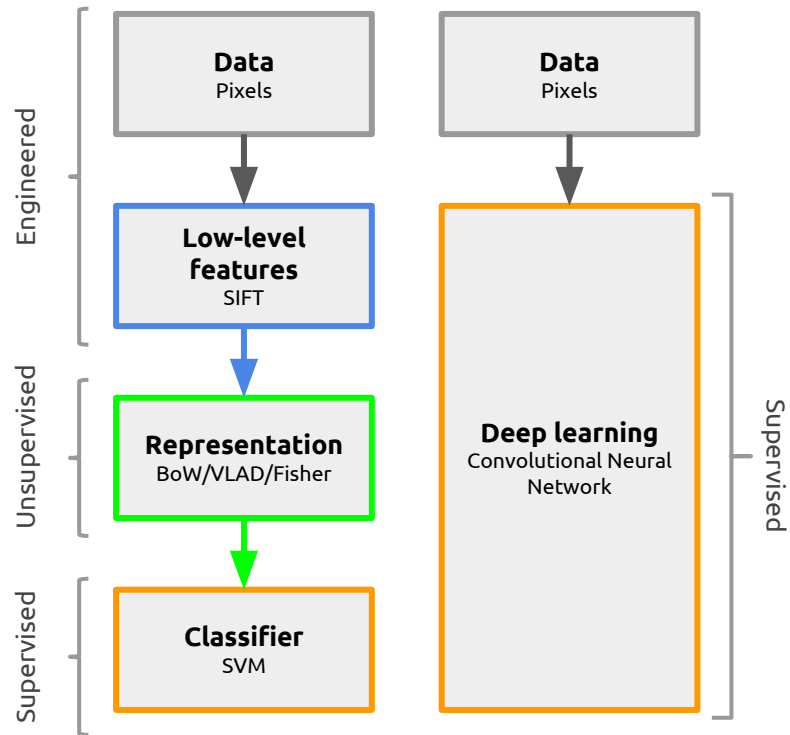
End2end model



Model maps raw inputs to raw outputs, no intermediate blocks.

Classic Machine Learning vs Deep Learning

- Old style machine learning:
 - Engineer features (by some unspecified method)
 - Create a representation (descriptor)
 - Train shallow classifier on representation
- Example:
 - SIFT features (engineered)
 - BoW representation (engineered + unsupervised learning)
 - SVM classifier (convex optimization)
- **Deep learning**
 - Learn layers of features, representation, and classifier in one go based on the data alone
 - Primary methodology: deep neural networks (non-convex)



Undergradese

What undergrads ask vs. what they're REALLY asking

"Is it going to be an open book exam?"

Translation: "I don't have to actually memorize anything, do I?"

"Hmm, what do you mean by that?"

Translation: "What's the answer so we can all go home."

"Are you going to have office hours today?"

Translation: "Can I do my homework in your office?"

"Can i get an extension?"

Translation: "Can you re-arrange your life around mine?"

"Is this going to be on the test?"

Translation: "Tell us what's going to be on the test."

"Is grading going to be curved?"

Translation: "Can I do a mediocre job and still get an A?"

