

# AIDL: PROJECTS



Session 1 (2019-03-05): Introduction

# Issey Masuda Mora

Education: Telecommunications Engineer



Currently working: Deep learning team lead at Vilynx



Technologies:



# Hand-raised poll

- Programming skills?
- Working on software development? Or management?

# Subject goals

**“Give them the tools to  
apply DL in the industry”**



# Expectations: what not to expect



# Expectations

- Implement NN models
- Ready to prod
- Useful for current job
- Product MVP
- Solve real-world problems
- State of the art

# Subject goals

**Product development**



**Ready to prod model**

**Implement DL model**



# Subject goals

## Tangibles:

- Implement DL models
- Optimize resource usage / increase efficiency
- Optimize models
- Deploy to production

## Intangibles:

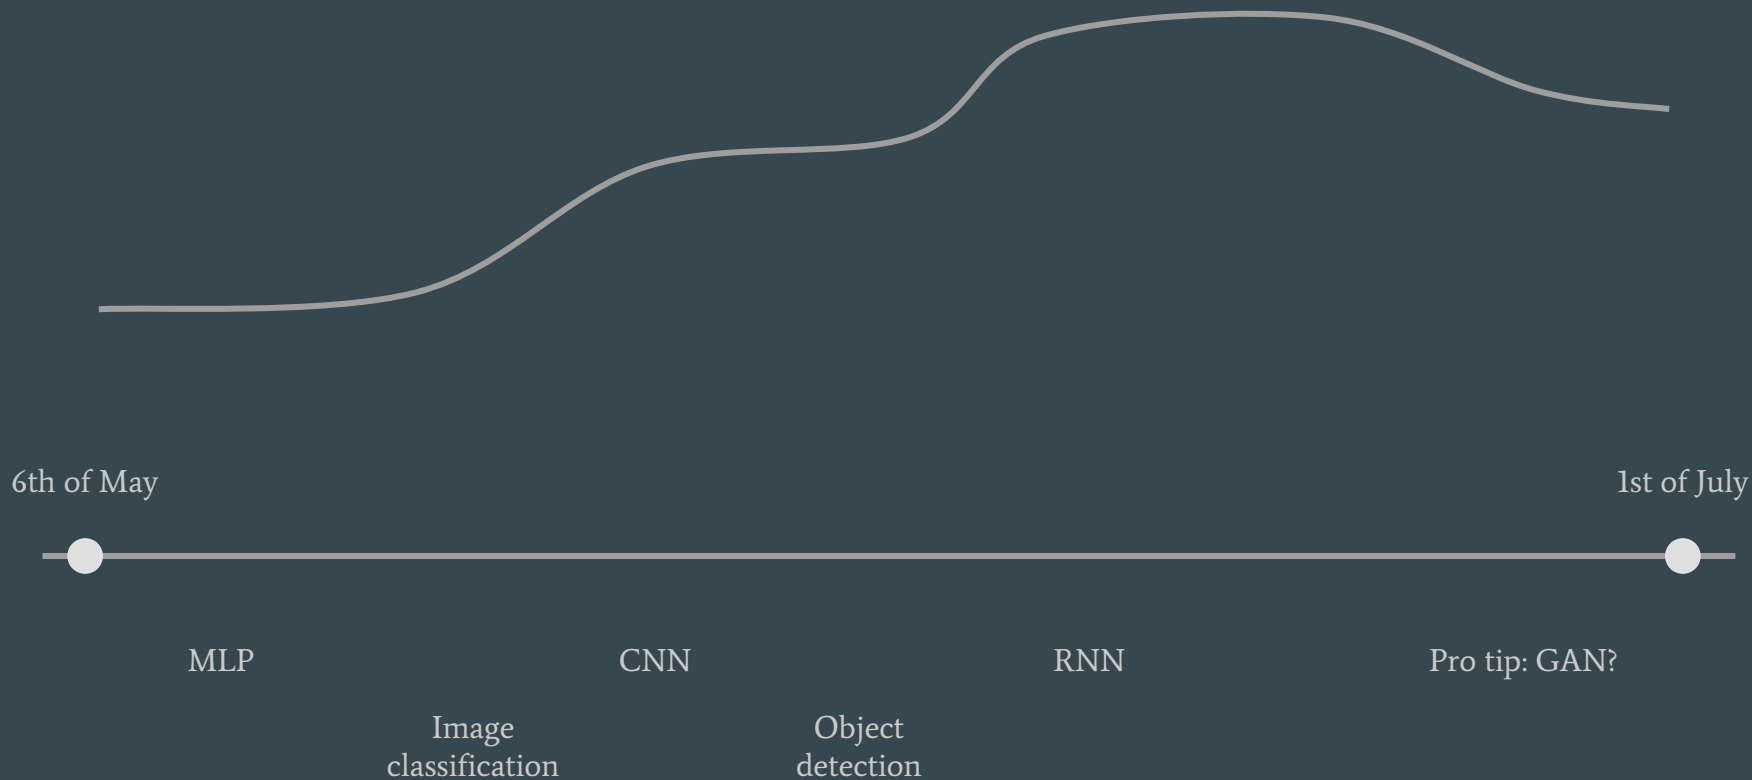
- Add a DL implementation into the roadmap
- Acquire a data scarcity mindset
- It's not all about DL!

# Course overview

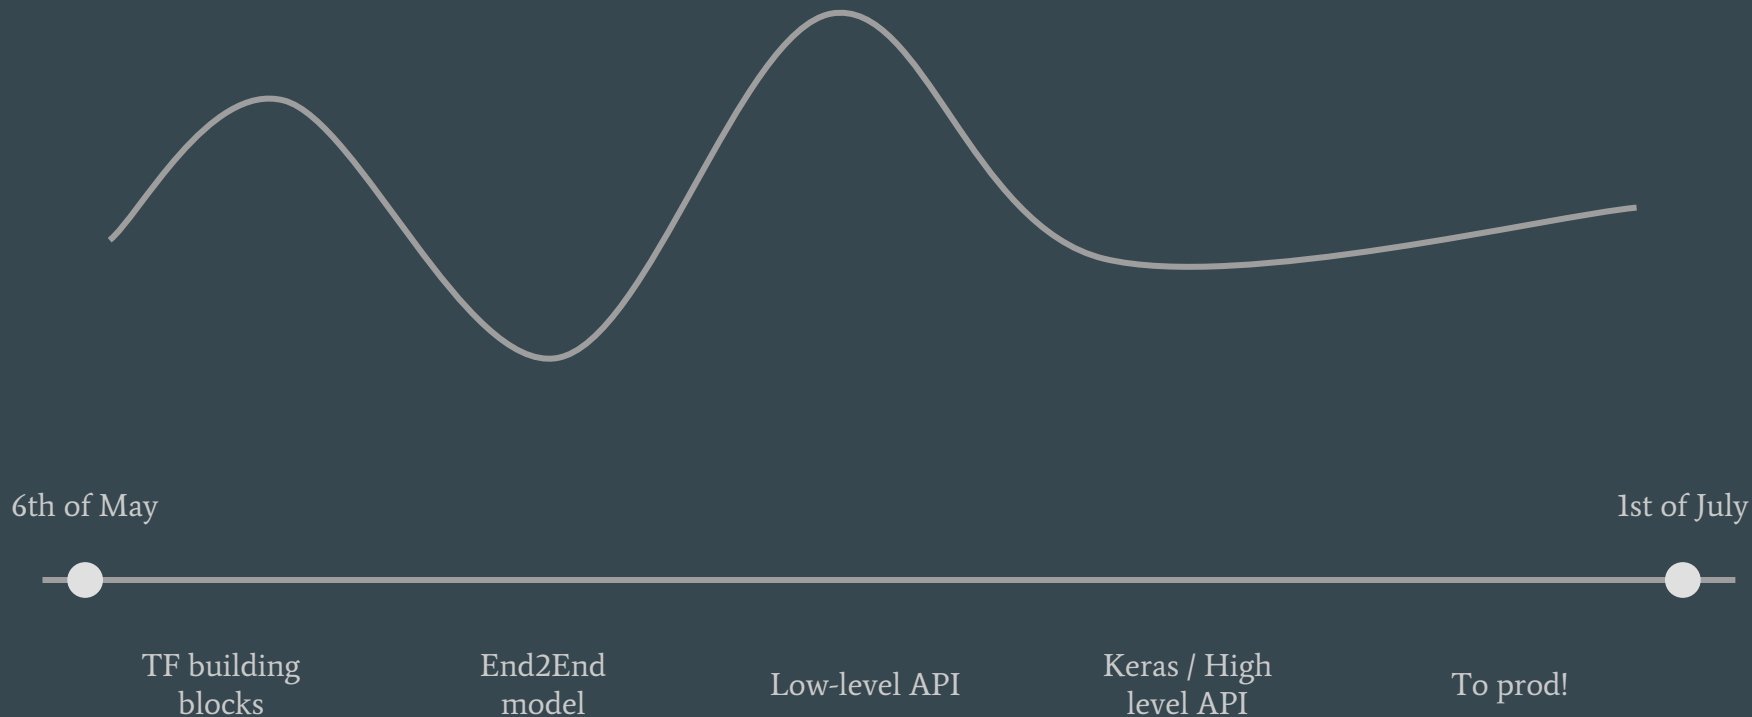
How is this all going to work?!?!?



# Deep learning



# Tensorflow



# Classes

# Classes

What they should be:

- Guide + DIY + review solution → what did we learn?
- Align each one with the subject goals

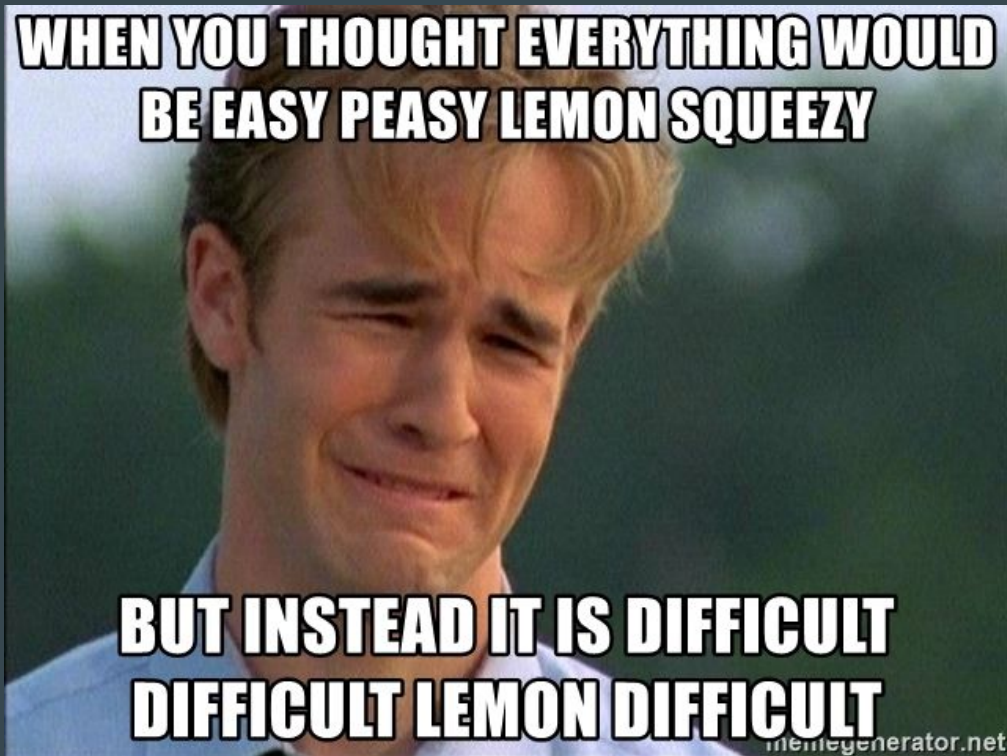
What they should not be:

- Medium blog post
- Tensorflow tutorial
- TF API reference summary

# Classes

~~Jupyter & Codalab~~

Old & good python





**Begin!**

# Computational graph

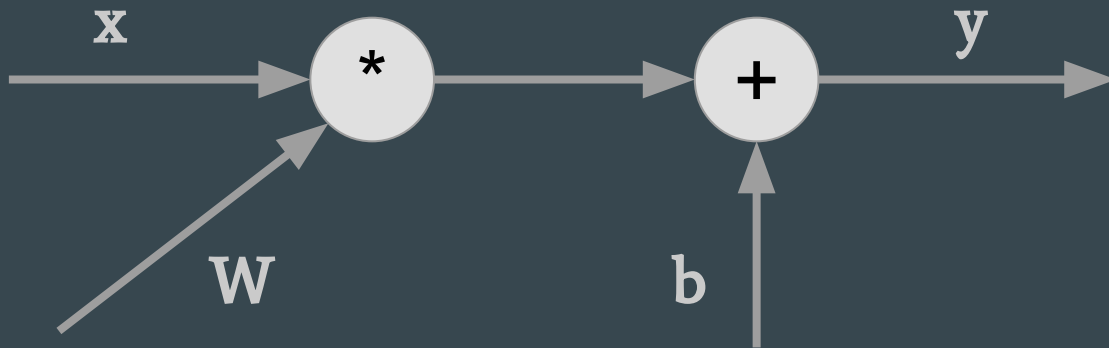
# What is it?

Nodes: operations

Edges: data / tensors

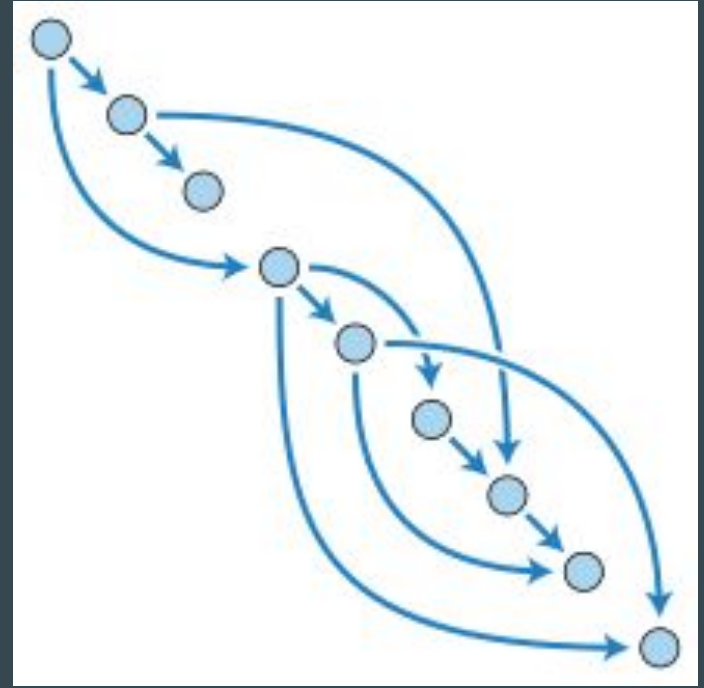
Example

*Linear regression:  $y = Wx + b$*



# Directed Acyclic Graph (DAG)

- Edges are directed from one node to another
- Nodes have a topological ordering
- No closed loops!



# Static computational graphs

Define-and-Run methodology.

Pros:

- Speed: can spend a long time optimizing the graph
- Memory: can predict and allocate all mem ahead of time

Cons:

- Inflexible: once compiled, it cannot be modified at run time
- Debugging: graph representation doesn't match code
- Static: more difficult to support flexible sized inputs

*Slide credit: Kevin McGuinness (adapted from the original one)*

# Dynamic computational graphs

Define-by-Run strategy

## Pros:

- Can change structure of NN at runtime (add layers, change shape, etc.)
- Support flexible sized inputs
- Easier to debug
- More natural to code

## Cons:

- Compile at runtime, can be slower
- Dynamic batching more difficult

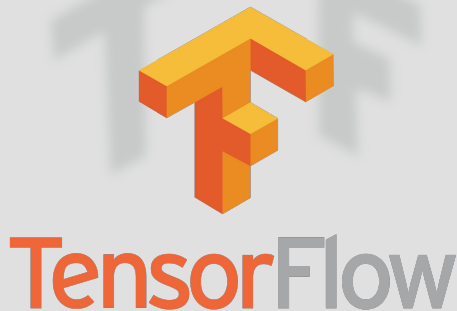
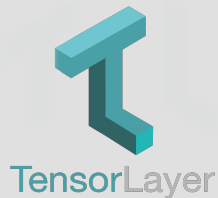
*Slide credit: Kevin McGuinness (adapted from the original one)*

# Deep learning frameworks



theano

DEEPLARNING4J



Lasagne

Caffe





# Framework basic comparison

## Tensorflow:

- Production ready
- Multiple languages supported (Python, C, Java, JavaScript, Go)
- Multi device (Android, iOS, Web browser)
- Mature with a big community
- Tensorboard
- Keras
- Static graph

... Tensorflow 2.0 coming soon! → dynamic graph

## Pytorch:

- Pythonic
- Easy to learn
- Framework rather than library
- Integration with Caffe2
- Dynamic graph

... Pytorch 1.0 finally released! → production ready

**Surrounding technologies**

# Git



- Course repository (for the coding sessions)
- Use a personal branch to keep track of your progress
- Master branch will be updated each week with the solutions of previous sessions & required material for the current one



# Virtualenv & virtualenvwrapper

- Isolate python environment

# Docker



- Development reproducibility
- Containerize production environment
- TF Serving



# NVIDIA

**CUDA:** “is a parallel computing platform and application programming interface (API) model created by Nvidia”



**cuDNN:** “is a GPU-accelerated library of primitives for deep neural networks. cuDNN provides highly tuned implementations for standard routines such as forward and backward convolution, pooling, normalization, and activation layers”

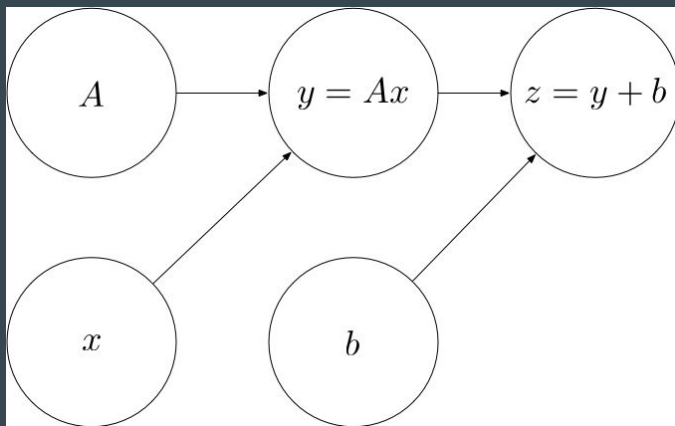


**TF preview**

# Computational graph & Dataflow

Programs are represented as directed graphs with data flowing through them where:

- **Nodes:** Operations of the program  $\rightarrow$  **tf.Operation**
- **Edges:** Data flowing through the graph  $\rightarrow$  **tf.Tensor**



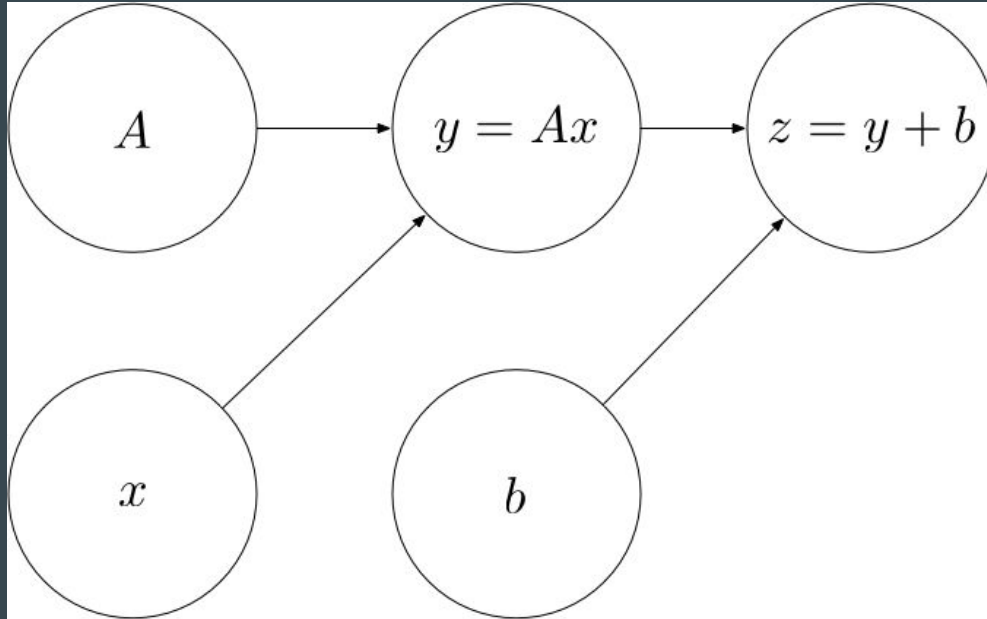


# From symbolic to real numbers

Two phases:

- **Definition phase:** build the computational graph → **tf.Graph**
- **Execution phase:** interact with the graph feeding data and fetching results.  
Run a subgraph of the original graph & change its state → **tf.Session**

# Micro example I: the graph



# Micro example II: the code

```
import tensorflow as tf
import random
```

```
graph = tf.Graph()
with graph.as_default():
    x = tf.placeholder(tf.float32, shape=[], name='x')
    A = tf.get_variable('A', shape=[], dtype=tf.float32,
initializer=tf.initializers.random_normal())
    b = tf.random_normal(shape=[], dtype=tf.float32, name='b')
    y = A * x
    z = y + b
```

```
with tf.Session(graph=graph) as sess:
    sess.run(tf.global_variables_initializer())
    result = sess.run(z, feed_dict={x: random.random()})
```

Fetching

Feeding

Definition phase

Execution phase

Demo time

# Questions?

